

Homework #2

Assignment For this second homework the exercise is related to the classification of handwritten digits of the MNIST dataset. The required tasks are the following:

- extract the data and labels from the given MATLAB file;
- propose a neural network (NN) for the classification of 10 digits;
- implement a cross-validation scheme and search the best hyper-parameters using a search procedure;
- calculate the mean classification accuracy;
- the output of the network should be of size ten plus a single integer number representing the predicted class.

In Figure 1 some examples of images taken from the dataset provided in class. The dataset does not need any preprocessing, so the images are used as-they-are for the training, validation and testing.

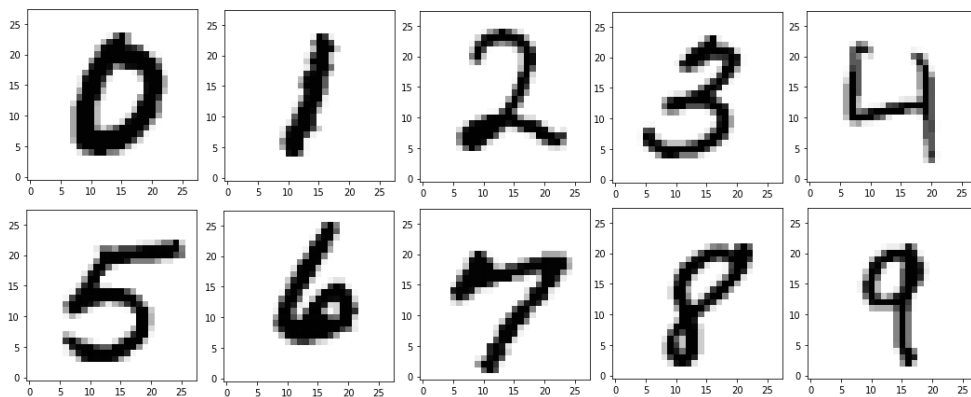


Figure 1: Some digits from the given MNIST dataset.

1) Model The proposed NN model is feed-forwarding, but different from the one seen in the laboratory: two *Linear* layers followed by a *ReLU* and a third *Linear* layer followed by a final *LogSoftMax*.

The input size (N_i) corresponds to the number of pixels per image, which is 784 (28×28), the hidden layers have a number of neurons (N_{h1} & N_{h2}) to be determined by the search procedure and the output size (N_o) is 10 as the number of classes to classify.

In Table 1 a scheme of the proposed feed-forward neural network. The model could also return as optional output just a single number for the chosen class.

Layer	Parameters
Linear	(Ni, Nh1)
ReLU	
Linear	(Nh1, Nh2)
ReLU	
Linear	(Nh2, No)
LogSoftMax	dim=1

Table 1: Scheme of the proposed model.

The aim of the algorithm is to classify ten different classes, so the final layer has to return ten outputs representing the single probability per digit: the highest probability represent the model prediction. In order to do so, I decided to use a *logarithmic Soft-Max* (LogSoftMax) layer which provides the probabilities for a input to be one of the possible classes.

As suggested in the *pytorch* documentation for classification problem concerning multiple classes, along with LogSoftMax is usually used a *negative log likelihood loss* (NLL-Loss), a loss function that gives a high value when the classification is unclear or wrong and a very low value if the prediction matches the expected value.

The chosen optimizer is the *stochastic gradient descent* (SGD) with a learning rate to be decided and a fixed momentum of 0.5.

2) Training From the original dataset made of 60 000 samples, the first 10 000 samples have been used as *Test* dataset and the remaining 50 000 as *Training* set. To speed up the process, I decided to use a *batch size* of 64.

While searching for the optimal parameters, only 10 000 samples have been used: they have been split again in *train* and *validation* set during a 3-fold cross validation.

In Table 2 the main parameters used to train the networks with 10 epochs before choosing the best ones.

Fixed parameters		Grid search	
<i>Ni</i> , number of inputs	784	<i>Nh1</i> , number of neurons	[32, 64, 128]
<i>No</i> , number of outputs	10	<i>Nh2</i> , number of neurons	[32, 64, 128]
<i>k-fold</i> , number of folds	3	Learning rate	[0.01, 0.001, 0.0001]

Table 2: On the left, the fixed parameters used to train the various networks; on the right, the parameters to try in order to obtain the best results.

The parameters providing the smallest loss in the *validation* dataset are saved and used to train the final network with 50 epochs. According to the *grid search* implemented, I obtained that the best choices are 64 hidden neurons in the first layer (Nh1), 128 for the second (Nh2) and a learning rate of 0.01.

Results With the aforementioned parameters and architecture, the obtained accuracy of the model is **0.974** on the *test* set.

In Figure 2 are reported some examples of images correctly classified a) & b) and wrongly classified c) & d).

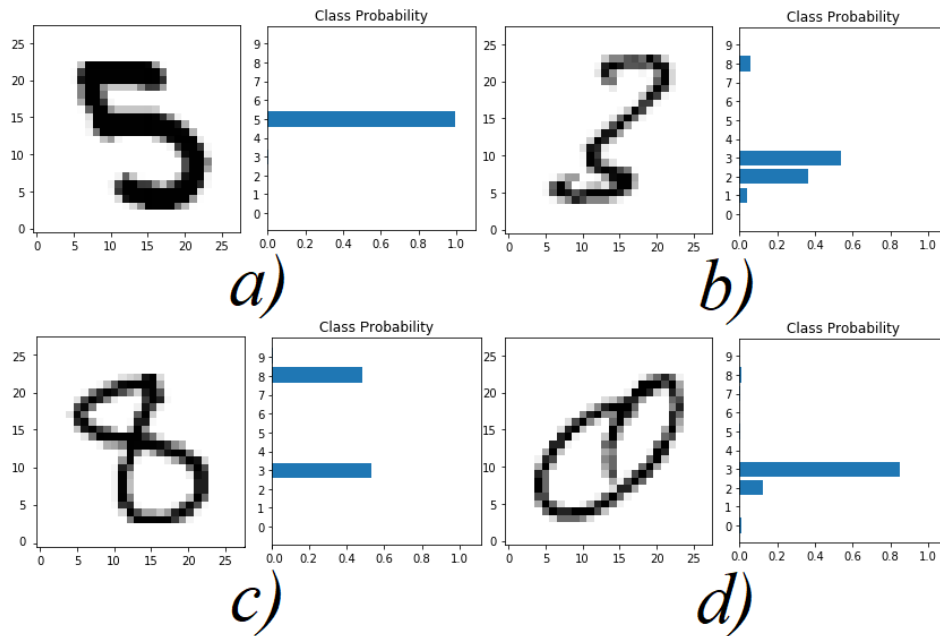


Figure 2: a) represent the majority of the results: the image is correctly classified without uncertainty; b) the model presents some ambiguity, but it's still able to correctly recognise the digit; c) the model fails to recognise the correct digit, but still recognise it as possible with relatively high probability; d) the model completely fail to recognise the digit.

Finally, in Figure 3, 4, 5 are plotted the receptive fields of the first, second and last layer respectively. They represent the area of the original image that could influence the activation of a neuron.




Figure 3: Ten of the receptive fields at the first hidden layer.




Figure 4: *Ten of the receptive fields at the second hidden layer.*




Figure 5: *Final receptive fields at the last layer of the model.*

External libraries:

- *scikit-learn*, for the k-fold;
- *scipy*, to import the .mat file.